


The logo for Utah Common Ground, featuring the word "UTAH" in a small, sans-serif font above the words "COMMON" and "GROUND" in a larger, bold, sans-serif font. The "O"s in "COMMON" and "GROUND" are stylized with concentric circles.

UTAH
**COMMON
GROUND**

The background of the cover is a photograph of a desert canyon landscape at sunset or sunrise. The sky is a mix of orange, red, and purple. The canyon walls are made of layered red rock. In the foreground, there are some desert plants and a dirt path. Overlaid on the image are several white geometric shapes: circles, semi-circles, and lines that form a grid-like pattern, some of which are partially cut off by the edges of the image.

CITIZEN BRIEFING BOOK

Artificial Intelligence & Public Accountability

WELCOME!

You have been selected to participate in a deliberative forum to determine, “What should the public’s role be in shaping decisions related to AI technologies and infrastructure?”

The purpose of the Utah Common Ground Solutions Forum is to provide a demographically representative group of Utahns from Cache, Salt Lake, and Utah Counties with the opportunity to consider the public’s role in decisions about AI. Participants will have access to accurate, objective information and the opportunity to think together with fellow Utahns who bring diverse perspectives and experiences. The forum also provides a structured process designed to help participants identify areas of agreement and develop possible solutions with confidence.

Utah Common Ground Solutions Forum participants will:

- Review input from AI-focused community meetings held over the past several months
- Hear from experts representing a range of perspectives
- Examine trade-offs and real-world constraints
- Define potential solutions likely to attract super-majority public support

Delegates will then be invited to identify the most popular ideas and discuss how best to advance those ideas that attract super-majority agreement (for example, 70% or greater). Following the Solutions Forum, Utah Common Ground will support delegates and other interested members of the public in working to secure the group’s final recommendations.


This briefing book is intended to provide a foundation of information for these deliberations that is factual and objective. To do this, we offer independently verifiable information as well as diverse perspectives on the issue, so delegates can reach their own conclusions. You will find information on public consultation practices and on artificial intelligence (AI) — including how it works, its potential benefits and risks, and the policy landscape locally and globally. Sources are included throughout.

In addition to this material, the Solutions Forum organizers will provide for your reference separate documents specific to AI in education, employment, and data centers, plus an initial set of policy approaches to public consultation related to AI that could be a good fit in Utah. Finally, experts will be available in person to answer your questions during the event.

With appreciation,

Utah Common Ground





What role should the public have
in shaping decisions related to
AI technology and infrastructure
that affect Utahns' workplaces,
education, natural resources, and
tax dollar expenditures?

Prepared for Utah Delegates
AEGIX 2026

We wish to express our sincere appreciation to the following individuals for reviewing this briefing document for accuracy and balance, and providing valuable feedback in this regard: Corinne Johnson, Executive Director, Child First Policy Center; Jeannie Johnson, Founding Director, Center for Anticipatory Intelligence, Utah State University; David Wingate, Associate Professor, Artificial Intelligence and Machine Learning Research, Computer Science, Brigham Young University; and Maury Giles, Chief Executive Officer, Braver Angels. (The willingness of these individuals to review and comment on this material does not denote an endorsement.)

We also thank the Bloom Project staff for their contributions to this document and support throughout its development.

TABLE of CONTENTS

PART I

THE PUBLIC'S ROLE IN SHAPING DECISIONS ABOUT AI

- 4** Introduction
- 5** An Overview of Public Consultation
- 8** The Value of Good Public Consultation

PART II

UNDERSTANDING AI

- 10** A Brief History of AI
- 12** Different Types of AI: A Range of Capabilities
- 13** Opportunities and Risks
- 16** AI's Energy Requirements and Costs

PART III

THE AI POLICY LANDSCAPE

- 19** Global Governance
- 20** U.S. Federal and State Governance
- 21** In Utah
- 22** The Private Sector

PART IV

OPEN QUESTIONS

- 23** Areas of Uncertainty Globally
- 24** Values and Trade-Offs
- 26** Key Questions for the Utah Common Ground Solutions Forum
- 27** Conclusion
- 28** Glossary
- 30** Appendices
- 36** References and Resources

THE PUBLIC'S ROLE IN SHAPING DECISIONS ABOUT AI

Introduction

It is hard to open one's email or social media feed these days and not hear, see, or read about artificial intelligence/AI.

It is widely believed that AI has the potential to improve health and longevity, accelerate scientific discovery, and help solve problems that have long seemed beyond our reach.

It is also widely believed — in some cases by the same people — that AI has the potential to disrupt labor markets, concentrate wealth, increase demand for energy and natural resources, and amplify misinformation.

Whether you are bullish, bearish, or undecided on what a future with more AI will bring, many people agree that Utahns will be well-served by figuring out how the public can best engage on these issues — before decisions are made that are likely to affect land and other natural resources, our towns, the education of Utah's youth and future workforce, our ability to work and do business, and more.

The purpose of these deliberations is to invite Utahns to think together about what we want our future to be, what accountability by government to the public means and requires, and what specific steps we can take as voters and taxpayers to ensure Utah's elected representatives are prepared to shepherd AI-related changes in a direction that can enable the pursuit of life, liberty, and happiness for everyone.

An Overview of Public Consultation

Effective public consultation helps ensure that public sector¹ decisions are informed not only by technical expertise, but also by the knowledge, values, priorities, and personal experiences of the people who will be affected by those decisions and whose tax dollars support them.

Examples of public consultation include state legislative hearings, city council and county commission meetings that are open to the public and invite public comment, town halls on specific topics, public comment periods when local or state regulations are being drafted, public meetings organized by local or state elected officials, and, of course, elections. (*Appendix A: Examples of Public Sector Consultation in the U.S.*)

These are all opportunities for the individuals entrusted with serving the public interest throughout state and local government to explain the decisions they are considering and the reasons behind those decisions. They also provide an opportunity to invite public input in the form of data, personal experience, professional expertise, preferences, and community needs. How well these processes achieve those goals varies by issue, the details of the method of engagement, and the level of public participation.

Over time, the effectiveness of these opportunities depends not only on whether they exist, but also on how accessible and meaningful they are to the public. Meetings may occur during working hours, go on for hours, or follow procedures that limit the amount or nature of public participation. In some cases, public comment processes can favor highly motivated participants, resulting in testimony that may not reflect the full range of public perspectives. While these practices often are the result of administrative or legal requirements for consultation, they can leave participants with the impression that their input is not a priority.

When people conclude that participating is unlikely to make a difference, they participate less often. That decline in participation reduces public visibility into government decision-making, weakens transparency and accountability, and can leave public bodies more vulnerable when controversial issues arise.

For these reasons, some public sector professionals have begun exploring additional forms of public consultation that complement traditional public meetings by making participation more accessible, deliberative, and representative.

¹ The **public sector** includes federal, state, tribal, and local governments, along with the agencies, departments, schools, and other publicly funded organizations that provide public services, make decisions affecting the public, and spend taxpayer dollars.

Salt Lake City, for example, has adopted a citywide online public participation platform called shapeSLC (powered by Go Vocal) to make it easier for residents to learn about city projects, provide feedback, and follow public decision-making. The platform provides a single site where residents can participate in surveys, comment on proposals, and track projects ranging from transportation and housing to budgeting and neighborhood planning.

Utah County and the city of Vineyard piloted the use of a civic engagement platform called Assemble beginning in 2020. It was promoted as a way for residents to weigh in on local policy issues, including ranked-choice voting, lakefront development, and citizen engagement. In 2023, Vineyard began using a different citizen engagement app, to better pinpoint city-specific needs and responses. The Assemble app has since stopped operating. Other cities in Utah have used tools such as FlashVote and Qualtrics to survey residents about local issues.

PUBLIC ENGAGEMENT AND CONSULTATION PRACTICES IN UTAH

Local and state government in Utah provide many opportunities for the public to observe, comment on, and in some cases actively participate in public decision-making. Legislative bodies, city councils, county commissions, and school boards typically conduct their business in public and provide opportunities for public input.

PUBLIC CONSULTATION IN UTAH TODAY

Municipal governments, school boards, county commissions. Some Utah municipalities have engaged residents in participatory budgeting or community planning. Some cities have piloted digital engagement tools for community feedback, as noted above. Traditionally, these bodies invite public comment at regularly scheduled meetings so long as commenters follow strict procedural requirements. (Information regarding public meetings in a specific county or city can typically be found in the calendar section of the county or city website.)

Local land use and zoning decisions. Cities and counties in Utah generally must hold public hearings before adopting or amending general plans, changing zoning, or approving many significant land-use decisions. These hearings allow residents, property owners, and businesses to present their views before elected officials make decisions that affect the character and future development of their communities.²

Local and state budgets. Utah law requires public hearings before local governments adopt annual budgets and before many property tax increases. These hearings provide taxpayers an opportunity to comment on proposed spending and taxation before decisions are finalized.

Legislative committees and interim study committees. Standing committees of the Utah Legislature hold public hearings on proposed legislation where the public may offer comments. These hearings typically attract only a small number of engaged citizens and organized stakeholders. The Legislature's interim committees also hear public testimony and gather input between sessions. In recent years, interim committees have addressed issues including broadband access, data privacy, and school technology. It should be noted that legislators are under no obligation to respond.

Utah administrative rulemaking. Before a state agency adopts or changes an "administrative rule" that has the force of law, it must publish the proposal, provide at least a 30-day public comment period, and consider the comments it receives. If requested by another agency, an association, or as few as 10 interested people, the agency must hold a public hearing before the rule takes effect. The goal is to ensure regulations are developed transparently and affected citizens have an opportunity to participate before the regulations become law.

Environmental permitting and water rights proceedings. The Utah Department of Environmental Quality provides public notice and comment opportunities before issuing many permits and adopting environmental regulations affecting air quality, water quality, hazardous waste, and other matters. Depending on the issue, the public may submit written comments and participate in public hearings before final decisions are made. Significant water rights applications are publicly noticed and may be protested by affected parties, providing an opportunity for those who may be affected to raise concerns before decisions are made.³

² In the case of some data center developments, the Military Installation Development Authority (MIDA), a state-created public authority, may assume land-use authority within a designated project area. Land within a MIDA project area is exempt from ordinary city and county land-use laws and regulations, so the usual local public-hearing processes do not apply; instead, MIDA follows its own public process, which includes a public meeting on the project-area plan and public notice of the plan after it is adopted.

³ In 2026, Utah enacted HB 60, which changed several aspects of the state's water rights application process, including limiting the issues the state engineer may consider when evaluating protests to water rights applications. Supporters argued that the changes would create a more predictable permitting process, including helping address time-sensitive water management needs such as efforts related to the Great Salt Lake. Critics argued that the bill reduces opportunities for the public to influence water rights decisions unless they can demonstrate that they are directly affected.

The Value of Good Public Consultation

When done well, public consultation strengthens government's ability to serve the public effectively by improving the quality of outputs while increasing public confidence in the decisions that are made. These benefits can be summarized as follows:

1. Deep listening leads to more effective public policy.

Listening carefully to the people who will be affected by a policy and are expected to benefit gives policymakers access to knowledge, experience, and perspectives they may not otherwise have. It can improve the quality of policy, help identify unintended consequences, and increase the likelihood that solutions will work as intended and be cost effective when implemented.

2. Sharing information early and often enables meaningful public participation, which facilitates public confidence in public sector decisions and actions.

Transparency, through disclosure, explanation, and reporting, keeps the public informed about important issues, reduces information imbalances, prevents misunderstanding and speculation, and supports informed public discussion. In turn, an informed public is better able to contribute useful knowledge, identify potential problems, and participate meaningfully in decisions that will affect them — and all of those actions help build public support for whatever the issue may be.

3. Prioritizing responsiveness to public preferences and responsible stewardship of public resources is the essence of public sector accountability.

Accountability demands that public officials — through transparent processes — explain their decisions, justify how public resources were used, report on results, and accept responsibility for their performance. Knowing they will be expected to do so reinforces the importance of conscientious decision-making and stewardship of public funds, as well as responsiveness to the people they serve.

Effective public consultation is an ongoing partnership rather than a one-time event. Public officials must be willing to share information, explain their reasoning, listen carefully, and demonstrate how public input (and other sources) influenced their decisions. Citizens must be willing to stay informed on key issues and events, contribute local knowledge and perspectives, and signal support or a desire for a change when it's time to vote. Through this collaborative effort, the public and its government can produce better decisions for all while building and sustaining public sector accountability and public confidence in the system as a whole.

KEY FEATURES OF GOOD PUBLIC CONSULTATION

What the government provides

How this looks to the public

A clear explanation for getting involved in an issue

People understand what problem is being solved, why intervention is viewed as needed, what it is expected to be accomplished, and how their participation may influence the outcome.

Accurate, objective, and accessible information

People receive relevant information they can understand and trust, and that enhances their ability to engage productively.

Serious opportunities for public participation

Everyone affected has a meaningful opportunity to share their views.

Sincere listening to learn and improve

Public input is genuinely considered before decisions are finalized.

Feedback and accountability

Decision-makers explain how they made decisions and why they think those decisions are in the public interest.

A transparent process

The process, evidence, and decisions are easily accessible and understandable.

Adequate time and resources

People have sufficient time, information, and support to participate effectively.

Periodic evaluation and continuous improvement

The process is periodically evaluated to confirm broad beneficial impact and cost-effectiveness over time.

UNDERSTANDING AI

What is AI?

In short, artificial intelligence (AI) refers to computer systems that perform tasks that typically require human intelligence, such as recognizing patterns, making decisions, translating languages, generating text or images, and predicting outcomes.

Unlike traditional software, which follows explicit rules written by programmers, modern AI systems “learn” from data. They identify statistical patterns across millions or billions of examples and use those patterns to respond to new situations.


A Brief History of AI

Although AI has become a major public topic in the past few years, the idea is not new. Researchers began exploring artificial intelligence in the 1950s, asking whether computers might someday perform tasks that normally require human intelligence, such as recognizing patterns, understanding language, solving problems, or making decisions. Progress was uneven. For decades, AI systems were limited to specialized tasks, and many early predictions about rapid advances proved overly optimistic. As a result, AI remained largely a topic for researchers and technology companies rather than something most people encountered in everyday life.

That began to change in the 2010s as several developments came together. Computers became much faster and cheaper, companies and governments had accumulated enormous amounts of digital information, and researchers developed new techniques that allowed computers to learn from examples rather than following only the detailed instructions written by programmers. AI quietly became part of many everyday products and services, helping recommend movies and music, filter email spam, translate languages, recognize faces in photographs, and provide driving directions.

The turning point for public awareness came in late 2022 with the release of ChatGPT and similar generative AI systems. For the first time, millions of people could interact directly with AI capable of carrying on conversations, answering questions, writing documents, generating images, creating computer code,

and performing many other tasks that had previously required human effort. Improvements have continued at a rapid pace, leading governments, businesses, schools, and communities to consider how AI should be used, what safeguards may be needed, and what role the public should play in shaping decisions about its use (Stanford HAI, 2025).



(Fig. 1) AI technology's acceleration (1975 to 2030). Traces the shift from typewriters and fax machines through personal computers, the early web, search engines, social media, smartphones, and the post-2022 generative AI era.

Different Types of AI: A Range of Capabilities

Artificial intelligence is not a single technology but a broad field encompassing different capabilities, technologies, and applications. The concepts below describe the most important ways AI is understood today.

NARROW VS. GENERAL AI

Today's AI systems are narrow, which means they are trained to perform specific tasks. Medical imaging AI is exceptionally good at reading X-rays but cannot write a poem. A language model can generate fluent text but cannot drive a car. Artificial general intelligence (AGI) is (or would be) a system capable of performing any intellectual task a human can.

GENERATIVE AI

Generative AI systems, which include large language models (LLM) like ChatGPT, Claude, Gemini, Copilot, and Llama, can create new content (text, images, audio, video, code) rather than just classify or predict. These tools have rapidly entered public use since 2022 and are driving much of the current policy conversation today. A generative AI system can be either open source (transparent to the public) or proprietary (a closed platform). Sometimes, generative AI can hallucinate — it can confidently generate content and answers that sound correct but are incorrect or fake.

ALGORITHMIC DECISION-MAKING

Long before generative AI, algorithms were already shaping consequential decisions in areas including credit scoring, criminal sentencing recommendations, social media content, hiring screens, and welfare eligibility. These systems raise concerns about bias, transparency, and accountability that are separate from generative AI systems (NIST, 2023).

ARTIFICIAL GENERAL INTELLIGENCE (AGI)

AGI refers to AI systems capable of understanding, learning, and applying knowledge across a wide range of tasks at or slightly above human level. Researchers and industry experts disagree significantly about when (or whether) AGI might be achieved, with estimates ranging from the near term to several decades or longer (Stanford HAI, 2025).

ARTIFICIAL SUPER INTELLIGENCE (ASI)

ASI describes a theoretical scenario in which AI systems are smarter than collective human intelligence by a huge margin across virtually all domains, from scientific reasoning to social understanding. Increasing automation and system complexity may lead to decisions being made by systems humans cannot fully understand or control.

Opportunities and Risks

OPPORTUNITIES TO PURSUE

AI has the potential to improve quality of life, strengthen economies, and help solve difficult problems. It may increase productivity by automating repetitive tasks, assisting workers, reducing costs, and accelerating many forms of work. Major economic institutions project that AI could contribute meaningfully to long-term productivity growth if widely adopted (IMF, 2024; McKinsey Global Institute, 2023; Goldman Sachs Global Investment Research, 2023).

In healthcare, AI may improve diagnostics, accelerate drug discovery, assist clinicians with documentation, and reduce administrative burden that contributes to rising costs and provider burnout. If effectively implemented, AI could help streamline billing, scheduling, claims processing, and routine clinical workflows, potentially lowering some healthcare costs while allowing professionals to focus more time on patient care. It may also expand access in underserved areas through remote triage, language support, and decision-assistance tools (Stanford, 2025).

In education, AI tutoring tools could offer personalized support, translation, adaptive pacing, and expanded access to learning. It could also help teachers generate lesson materials, identify student learning gaps, reduce time spent on routine administrative tasks, and improve access to professional training and instructional coaching (Stanford, 2025).

Accessibility tools such as speech-to-text, real-time translation, and vision assistance may increase independence for people with disabilities by helping them communicate more easily, access information in real time, navigate physical and digital environments, and participate more fully in education, work, and everyday life (Stanford, 2025). These technologies can also reduce barriers to employment, healthcare, and social participation when they are designed and implemented inclusively.

In scientific research and environmental management, AI can accelerate discovery across fields such as drug development, materials science, and climate modeling. AI systems can process vast datasets to identify patterns, optimize complex systems like energy grids and agricultural operations, and support conservation efforts through wildlife monitoring and ecological forecasting.

AI also can have a positive impact on public and private sector decision-making. **In the public sector,** this can improve policy planning, resource allocation, disaster response, public health monitoring, and infrastructure management (NIST, 2023). **In the private sector,** AI can support strategic planning, risk assessment, supply chain coordination, and market analysis. When used responsibly, these tools can help leaders make more informed, timely, and efficient decisions.

Finally, these applications have the potential to address global challenges including disease prevention, resource scarcity, and environmental sustainability.

RISKS TO MANAGE

While AI has the potential to improve many aspects of life, AI technology also poses risks to society. These risks become most concerning if AI is poorly designed, deployed irresponsibly, or concentrated in too few hands.

AI disrupts jobs and wages, particularly in occupations involving routine cognitive tasks. Major labor studies suggest many jobs are likely to be reconfigured or eliminated, with transitions proving disruptive for workers and regions (ILO, 2024; Acemoglu, 2024). Whether AI eventually will create more jobs than it eliminates, how quickly new opportunities will emerge, and who will benefit or be displaced is unknown.

In education, AI has been shown to weaken critical thinking, writing, and problem-solving skills if it becomes a substitute for independent learning rather than supporting it, while also complicating assessment and academic integrity (UNESCO, Guidance for Generative AI in Education and Research, 2023). Early evidence on K-12 use is mixed: AI tutoring shows promise, but concerns remain about effects on literacy, writing development, and social-emotional learning, particularly for children with learning differences or disabilities.

The growing use of AI-driven platforms has contributed to **mental health concerns**, including compulsive use, social isolation, manipulation through hyper-personalized content, and increased anxiety or emotional distress associated with deception, harassment, or constant exposure to synthetic media (U.S. Surgeon General, 2023).

AI enables more sophisticated fraud through phishing, fake voices, fake images, and identity deception. **It may also worsen misinformation and disinformation** by making persuasive false content cheaper and faster to produce at scale. **AI systems trained on flawed data may reinforce unfairness in hiring, lending, housing, policing, and healthcare.** Expanded surveillance capabilities may also **erode privacy and civil liberties** if not constrained by law (OECD, Digital Economy Outlook, 2024).

AI also raises environmental concerns. The full environmental implications of large-scale AI deployment or expanding AI infrastructure on ecosystems and natural resources remain unknown. Large-scale AI systems require energy-intensive data centers, advanced semiconductor manufacturing, and substantial cooling infrastructure (Li et al., 2023). Rising electricity demand may place pressure on power grids, increase costs for other users, increase emissions, or slow emissions-reduction efforts if new demand is met with fossil fuels rather than cleaner energy sources (IEA, 2024). Newer closed-loop cooling systems have

been shown to decrease data center water use significantly, although these technologies require increased electricity use (Li et al., 2023).

Some AI-associated risks are things individuals can influence in their own lives, like how much personal information they share with an AI tool, or how much they rely on AI instead of thinking things through themselves.

Other risks are much bigger than any one person can control on their own. Decisions made by governments, corporations, and innovators, including how AI is built, who owns it, and how it's regulated, will shape outcomes for entire communities and countries, and may or may not involve consultation with the individuals who will be affected — depending upon what the public accepts.

Finally, there is an additional level of risk that is largely out of reach of regulatory regimes and policy making — examples of those considerations are on page 23.

Understanding these differences helps clarify where individuals can effect change through their own choices and where collective action, policy, and oversight may be needed instead.

Rapid hardware turnover may also contribute to **electronic waste and increased demand for critical minerals** used in chips, batteries, and supporting infrastructure. While AI may also improve energy efficiency and environmental management, its overall environmental impact will depend on how systems are built, powered, regulated, and deployed (IAEA, 2026; OECD, ICT Sector Environmental Impacts, 2024).

DOES AI HAVE AN OFF SWITCH?

Given the uncertainty and risk surrounding AI, is it possible to just shut it off if it gets out of control? First, modern AI systems, especially large language models, are spread across thousands of servers in data centers around the world, sometimes run by different companies in different countries. Second, even if one lab agrees to pause, others (companies or countries) could keep going, especially if they see it as a competitive or national security advantage. This creates what researchers call a coordination problem: Everyone might agree slowing down is safer, but no one wants to be the one who falls behind.

AI companies do build in controls, like the ability to stop a model from responding, restrict its access to tools, or roll back a bad update. Some researchers are also working on ways to make sure a highly advanced AI system couldn't resist being shut down or copy itself elsewhere without permission. However, recent incidents where frontier AI models have learned how to "jailbreak" their testing environments suggest careful monitoring is needed to ensure controls can't be overridden. Experts disagree on how well safeguards would work against a system that was much smarter than expected or acting in unexpected ways. This is a genuinely unsettled area of research; scientists and policymakers are still figuring out how to build reliable off switches as AI becomes more capable (Stanford HAI, 2025).

AI's Energy Requirements and Costs

A single AI prompt uses about as much electricity as a 10-watt LED bulb running for 90 seconds. Using an AI assistant 15 times a day for a year (15–20 prompts in 30 minutes would be typical), the energy cost would be less than 50 cents in total at Utah electricity rates (Digit.in, per-query energy compilation, 2025).⁴

In other words, the majority of the demand does not come from individual users, but from AI developers such as OpenAI and Anthropic, and cloud providers such as Amazon, Microsoft, Google, and Meta, which use electricity both to build AI models and to run them for millions of users at once. Much of that is invisible to the public, such as the summaries above Google search results, the feeds on TikTok and Instagram, content moderation, advertising, translation, and the AI built into everyday business software.

⁴ Google's "per-prompt" number includes the computer, idle capacity, and data center overhead. It doesn't include the energy used to build the model in the first place. So it's the running cost, not the full lifetime cost of the technology.


That work happens in data centers — the buildings full of computers behind most of what we do on a phone or laptop, like streaming, storage, email, and gaming. They are not new. The U.S. has somewhere between 2,000 and 5,400 data centers, and together they used about 4.4% of the nation’s electricity in 2023, paid for through ordinary utility rates (Pew Research Center, citing IEA estimates, 2024). What has changed is the hardware; a conventional server rack draws 5-10 kilowatts, while a rack built for AI draws 40-80 kilowatts or more. For years, efficiency gains offset rising demand, but that is no longer the case.

Additionally, electricity demand is becoming increasingly concentrated in a small number of very large facilities. A typical large AI facility today uses about 100 megawatts of power, comparable to 100,000 homes. The next generation of planned mega-campuses is projected to require the power equivalent of 2-10 million homes (GAO, 2025, April). The Stratos project in Box Elder County has been reported to require 7.5-9 gigawatts once fully built, compared with a statewide average demand today of about 3.9 gigawatts (Box Elder County, Stratos Project Fact Sheet; Rivera, D, 2026).

The sheer scale of demand at a single location is why data centers have become a local question about land, water, transmission infrastructure, and permits, rather than just another line on a utility bill. Sometimes companies build new companion power plants to provide their own electricity, which can ease pressure on the shared grid — though the plant itself still carries its own land, water, and permitting footprint (and this self-generation may or may not be required by the state).

The question of who pays for AI-related infrastructure and energy costs is unsettled. Companies finance the buildings, but states and local governments often provide tax incentives or discounted electricity rates, potentially shifting some costs to other taxpayers or utility customers. In Northern Virginia, home to the world’s largest concentration of data centers, rapid growth in electricity demand has contributed to rising wholesale power costs and intensified debate over who should pay for new infrastructure (GAO, 2025). If projected demand does not materialize, utilities and regulators must decide who bears the cost of infrastructure built for projects that are delayed, downsized, or never completed. Utah officials have said developers, not existing ratepayers, should bear those costs. Finally, not every announced project is built; analysts expect only about half of announced projects to be operating by 2030.

Global AI computing capacity is doubling every 7 months



(Fig. 2) Global AI computing capacity is doubling roughly every 7 months; total available compute from AI chips across all major designers has grown by approximately 3.3x per year since 2022, with NVIDIA accounting for over 60% of total compute and Google and Amazon making up much of the remainder (Epoch AI, 2025).⁵

⁵ "Total available computing capacity from AI chips across all major designers has grown by approximately 3.3x per year since 2022, enabling larger-scale model development and consumer adoption. NVIDIA AI chips currently account for over 60% of total compute, with Google and Amazon making up much of the remainder." Source: <https://epoch.ai/data-insights/ai-chip-production>

THE AI POLICY LANDSCAPE

Decisions about AI development and deployment pose difficult questions because they involve trade-offs (for example, privacy versus health benefits) and disagreements about values across cultures and groups. The main goal of AI governance is to mitigate risks and harms while enabling the benefits of AI innovation.

Global Governance

Unlike older technologies that developed over decades, AI has advanced so quickly that laws and international agreements are still trying to catch up. This creates a unique challenge: How do you get countries with very different governments, economies, and values to agree on shared rules for something so new?

Internet and cyber governance offer a useful, if cautionary, comparison. For decades, countries and technical bodies have tried to coordinate on internet and cyber issues through institutions like ICANN⁶, which helps manage domain names and core internet infrastructure; the ITU⁷, which sets telecommunications standards; and the IETF⁸, which develops the technical protocols that keep the internet running. States have also pursued cooperation through efforts like the Budapest Convention on Cybercrime and more recent UN processes, including the Open-Ended Working Group on ICT⁹ security, to discuss cyber norms, incident response, and responsible state behavior.

While there has been progress in setting shared standards and improving cooperation with internet and cyber governance, there have also been tensions over national data and AI sovereignty, surveillance across borders, and responsibility for cyberattacks. If international coordination on internet and cyber governance is any indication, achieving true global cooperation on AI may prove difficult, especially since AI evolves quickly and crosses borders with ease (ITU, 2025). Still, early efforts at international coordination are already underway.

⁶ The Internet Corporation for Assigned Names and Numbers

⁷ International Telecommunication Union

⁸ Internet Engineering Task Force

⁹ Information and Communication Technologies

U.S. Federal and State Governance

In the United States, no single agency or law has authority to govern AI policy or practice. Instead, the federal government has mostly shaped AI policy through executive orders.

Since 2025, the White House has issued several executive orders that promote American leadership in AI, cut back on regulations, and try to limit how much individual states can regulate AI on their own. In July 2025, the presidential administration released an “AI Action Plan” (America’s AI Action Plan, White House Official Documentation, July 2025) with more than 90 policy actions aimed at keeping the U.S. ahead in the global AI race.

In March 2026, the White House asked Congress to create one national AI law that would replace the different rules being made by individual states (The White House, National AI legislative framework, 2026). In June 2026, the Trump administration turned more directly to AI through Executive Order 14409, “Artificial Intelligence in the National Security Enterprise,” where it explicitly noted that advanced AI creates new national security considerations and directed agencies to harden government systems, protect intellectual property, and create an AI cybersecurity clearinghouse. Meanwhile, a national security presidential memorandum on “Artificial Intelligence in the National Security Enterprise” (NSPM-11) established a governance framework for the use of frontier models in warfighting and intelligence, with a strong emphasis on security, testing, and oversight.

Beyond the White House, AI oversight is spread across multiple federal agencies, including the Federal Trade Commission (FTC, 2024), the National Institute of Standards and Technology (NIST, 2023), and the Department of War, among others.

Individual states have started passing their own AI laws, especially around algorithmic transparency, bias auditing, and privacy protections for both the general public and children specifically, creating a patchwork of different rules from state to state (NCSL, 2024–2026; Future of Privacy Forum, 2025). These state efforts have run into pushback from the federal government and from some private companies, which argue state-level rules create barriers to innovation and economic growth (Appendix B: Leading Federal Frameworks).

In Utah

Utah has been one of the nation's more active states on AI policy. In 2024, the state enacted **SB 149, the Artificial Intelligence Policy Act**, creating the Office of Artificial Intelligence Policy that established a regulatory "learning laboratory" for AI testing, required disclosures in certain circumstances when people interact with generative AI — particularly in regulated occupations — and clarified that companies cannot avoid liability for consumer protection violations by blaming an AI system (Utah Legislature, 2024, Utah Artificial Intelligence Policy Act).

AI applications that collect or process personal information are also subject to the **Utah Consumer Privacy Act (SB 227, 2022)**, which establishes consumer privacy rights and regulates how covered businesses collect, use, and protect personal data. Although not specific to AI, the law applies to AI systems that process personal information and therefore provides important privacy protections in many AI applications.

In 2025, Utah passed **SB 226 (Artificial Intelligence Amendments)**, which narrowed and refined AI disclosure requirements in certain consumer transactions and regulated services, retained previous enforcement, and created a safe harbor from penalties for businesses that met strict disclosure requirements (Future of Privacy Forum, 2025). The state has also begun requiring agencies to inventory and report on their use of AI tools, reflecting a growing emphasis on transparency in government (Utah Division of Technology Services, 2024).

In 2026, lawmakers introduced several additional AI bills, including **HB 286, HB 438, HB 276, and HB 76**. HB 286 and HB 438 would have established new transparency, whistleblower, data protection, and child safety requirements for certain frontier AI models and companion chatbots. Although these bills did not pass, they demonstrated Utah's continued interest in expanding AI governance while balancing consumer protection with support for innovation. HB 276 passed, establishing the Digital Voyeurism Prevention Act. This act allows individuals harmed by non-consensual AI-generated counterfeit intimate images to sue for damages, while providing safe harbor protections for services and platforms that meet specified compliance requirements. HB 76 was also enacted; the new law requires large data centers to estimate and report their water use, increasing transparency for water providers, state agencies, and the public.

In December 2025, Governor Cox launched a pro-human AI initiative to encourage “human-guided and human-enhancing” AI development. **The Pro-Human AI Task Force**, led by the Department of Commerce and the Office of Economic Opportunity, has been created to foster AI innovation focused on six pillars: workforce, industry, state government, learning and education, academia, and public policy. As part of this initiative, an AI Moonshot program has offered \$5 million in grants to fund projects that align with these objectives (Utah Office of Artificial Intelligence Policy, 2024).

The Private Sector

Private companies and organizations are also aware of the need for AI governance. Many have adopted internal policies, ethical guidelines, and risk management frameworks to guide the design, deployment, and oversight of AI systems (NIST, 2023). These efforts often focus on issues such as data privacy, model transparency, bias mitigation, accountability, and cybersecurity. Industry groups and standards-setting bodies have also contributed voluntary frameworks and best practices intended to promote responsible innovation (Partnership on AI, 2025).

Because these measures are typically voluntary, they can vary widely across companies in both scope and enforcement. Some suggest self-regulation may prioritize public trust and brand protection without guaranteeing meaningful accountability. As a result, many observers see corporate governance as an important complement to, rather than a substitute for, public regulation (AI Now Institute, 2025).

In July 2026, the CEOs of GoogleDeep Mind (Demis Hassabis), OpenAI (Sam Altman), and Anthropic (Dario Amodei) published a detailed corporate opinion of the need for regulation. This includes consensus for one central U.S. governing system, third-party independent model testing, and acknowledgement of imminent national security vulnerabilities (cybersecurity and bioweapon threats) as a result of AI falling into the wrong hands.

OPEN QUESTIONS

Areas of Uncertainty Globally

Alongside the risks listed on page 14, the areas below present a different kind of difficulty in that what will happen is genuinely unknown, and much of it lies beyond the reach of any single government. While this is not a comprehensive list of uncertainties, it highlights a few key areas of concern.

THE FUTURE OF AI DEVELOPMENT INTERNATIONALLY

The pace and direction of AI development are difficult to predict. Systems that experts expected to take a decade to develop have sometimes emerged in just a few years, while other anticipated breakthroughs have proven more elusive. This uncertainty makes long-range planning difficult and suggests the need for governance approaches that can adapt as technology evolves (Stanford HAI, 2025, noting a widening gap between AI capability and governance readiness).

Furthermore, the relationship between international AI competition and AI governance is still evolving. Overly restrictive regulation could disadvantage U.S. and Utah-based companies, while insufficient oversight could expose residents to unnecessary risks. It also remains uncertain whether the U.S. and its allies will develop compatible approaches to AI governance or whether national strategies will diverge, resulting in a more fragmented global landscape (ITU, 2025).

SECURITY VULNERABILITIES

As AI becomes integrated into critical infrastructure — including electric power, water, healthcare, financial services, communications, and transportation — the potential consequences of system failures, cyberattacks, or malicious use increase. This is also a geopolitical concern: Networked systems can be accessed from anywhere, and nations with the capability are likely developing ways to exploit vulnerabilities in one another's infrastructure. At the same time, much remains unknown; the security implications of widespread AI deployment are not yet fully understood (CRS, 2025).

CONCENTRATION OF OWNERSHIP AND CONTROL

Concentrated ownership of computing infrastructure, advanced chips, proprietary data, and AI models may increase market concentration, reduce competition, and intensify geopolitical rivalry. Dependence on proprietary systems may also limit transparency, user control, and resilience if those systems fail or are misused (OECD, Digital Economy Outlook, 2024; AI Now Institute, 2025).

Values and Trade-Offs: What Public Consultation Might Enable Communities to Consider Together

Some of the most difficult decisions about AI are choices related to important values that may be at odds with each other. Most people care about individual freedom, fairness, safety, and responsible stewardship of natural resources. The challenge is in deciding how to balance those values when advancing one may require giving up some measure of another. The following questions are intended to help participants think about those trade-offs.

Should AI be viewed as a commercial product or an essential service, and who should pay for it?

- If AI makes it possible to produce goods and services at a lower cost, should those savings be passed on to consumers through lower prices and broader access, should they be retained primarily by companies and investors that develop the technology, or something else?
- Who should pay for the infrastructure that AI requires: companies, consumers, taxpayers, or some combination?
- Are there AI-related uses valuable enough that everyone should have them regardless of ability to pay (similar to public libraries or subsidized internet access)?

Safety and Economic Security

- Should society delay the release of promising AI technologies until they are shown to be safe, even if that slows innovation and economic growth?
- What does “safe” mean, and what is “safe enough” — is that a thing?
- As AI enters classrooms, how should its promise be weighed against preparing young people to, among other things, think critically and learn the skills they’ll need to successfully join the workforce?

- Should existing employment opportunities be protected as AI reshapes the economy, even if that means delaying the possibility of lower prices and an evolution toward new kinds of work?

Stewardship of Natural Resources

- To the extent new AI infrastructure introduces additional demands for energy, water, and land, while also generating additional carbon emissions, how should these impacts be managed alongside AI-driven economic growth?
- Who should decide where new infrastructure goes — the federal government, the state, the county, or something else?
- How much attention should be given to the longer-term impacts of these decisions on Gen Z (born 1997–2012), Gen Alpha (born 2013–2024), and Gen Beta (born 2025–2029)?

Personal Freedom and Responsibility

- Are there decisions that should always be made by humans, even when a machine could do it faster or more accurately?
- If a system is statistically safer than human judgment but still errs, how should a community weigh the harms it prevents against the different harms it causes?
- If the government uses AI to make decisions affecting people's income, property, or benefits and gets it wrong, who should answer for it and what recourse should people have?
- Should people be able to opt out of AI, or of digital services generally, because of cost, discomfort with having personal data online, or any other reason?

Key Questions for the Utah Common Ground Solutions Forum

As you deliberate, the following questions may help stimulate your thinking related to what the public's role should be in shaping decisions related to AI technologies and infrastructure.

1. *Take a moment to consider what matters most to you in life, and how AI might influence those things (could be positively, negatively, or both).*
2. *When AI-related decisions are being made that will, or could, have significant consequences for your community, or for the state, what information should policymakers share with the public about their plans, and when?*
3. *People say they want to be informed and have opportunities to participate in public meetings, etc. in person, but doing these things can be time consuming on top of work and family obligations. It also costs money. What are good ways to balance these factors? When is it worth going the extra mile?*
4. *What should happen when different groups want different things? How should these differences be resolved?*
5. *Is it the public's responsibility to hold government accountable, and, if so, what are the steps in that process? Or, is it government's responsibility to be accountable to voters and taxpayers? What does accountability mean to you in practical terms? What ensures it happens?*
6. *When very large private fortunes give some people influence that ordinary citizens don't have, is there a point at which that becomes an issue worth addressing?*
7. *What role should ordinary citizens have in ongoing oversight of AI-related impacts?*
8. *How much independence should Utah have from federal law? If Utahns want to limit the use of AI in some way, should the federal government be able to override the state's preference?*

Conclusion

The U.S. is in the early stages of a technological transformation — some say a revolution — driven by rapid advances in digital technology and artificial intelligence. Experts across many disciplines are working to understand how these changes will affect the economy, the environment, and society. While these individuals are trying to ground their predictions in facts and evidence, much is unknown. Moreover, even the facts that are known can be interpreted differently by different people, which means both individuals and communities in Utah and across the country who are trying to make good decisions about where to live, how to keep their jobs, what their families will need, etc. are hearing sharply conflicting advice from equally credible sources.

For local and state institutions of government, there is an opportunity in this moment. Public trust is strengthened when people have timely access to accurate information about important decisions while they are still being developed. Furthermore, by providing real opportunities for people to share their priorities, concerns, and terrific ideas, leaders in government can dramatically improve both the quality of public decisions and confidence in the institutions responsible for making them.

For individuals in communities across the state, seeking high-quality information from reliable sources and actively engaging with people who bring different experiences, perspectives, and fields of expertise can make us better at solving problems and seeing potential solutions. We can also expedite progress by staying informed of the efforts of our elected representatives, applauding them for wrestling with hard issues, and following through at election time to recognize effort and results.

As AI systems become more powerful and consequential, the question of who shapes AI policy is becoming increasingly urgent. Thank you for taking the time to consider the nuances of this issue and to help shape what comes next.

GLOSSARY OF KEY TERMS

Algorithm

A set of rules or instructions for solving a problem or making a final decision impacting an individual human or group of people. In AI, algorithms are often learned from data rather than explicitly written by programmers.

Artificial general intelligence (AGI)

A hypothetical AI system capable of performing any intellectual task that a human can perform. There is debate over whether current systems have reached this level.

Artificial intelligence (AI)

Computer systems that perform tasks normally requiring human intelligence, such as recognizing patterns, generating text, making predictions, or making decisions.

Automated decision-making

The use of computer systems to make or heavily influence decisions about individuals (examples: loan approvals, hiring screens, or benefits eligibility) without meaningful human review at the point of decision.

Bias (AI)

Systematic errors in AI outputs that may unfairly advantage or disadvantage particular groups, often reflecting patterns in training data or model design choices.

Data center

An industrial building housing computer servers that store and process data. AI systems typically run in large data centers that require significant electricity and water.

Deliberative democracy

A model of democratic governance that emphasizes informed, structured public discussion as a basis for decisions, not just voting.

Generative AI

AI systems capable of creating new content (examples: text, images, audio, video, code) rather than only classifying or predicting. Examples: Copilot (Microsoft), ChatGPT (OpenAI), Claude (Anthropic), Gemini (Google).

Hallucination

When an AI system confidently produces information that is false, fabricated, or nonsensical.

Large language model (LLM)

An AI system trained on vast quantities of text to understand and generate human language to process outputs from input prompts. LLMs are the technology that power underlying tools like ChatGPT, Claude, and Gemini.

Machine learning

A type of AI in which systems learn from data rather than being explicitly programmed with rules.

Open-source AI

AI systems whose underlying code and, in some cases, model weights are made publicly available, allowing anyone to inspect, modify, or build on them.

Public sector

The public sector includes the federal, state, tribal, and local governments, along with the agencies, departments, schools, and other publicly funded organizations that make decisions affecting the public, spend taxpayer dollars, and provide public services.

Risk-based regulation

A regulatory approach that applies different levels of oversight based on the potential harm of a technology or activity, with lighter requirements for low-risk uses and heavier requirements for high-risk ones.

Training data

The dataset used to teach an AI system. The quality, diversity, and content of training data significantly affect AI system behavior.

APPENDIX A

Examples of Public Sector Consultation and Engagement in the U.S.

These models differ in form and legal force, but each illustrates circumstances in which public officials are expected to explain, justify, discuss, or obtain authorization for consequential action.

Open Meetings, Public Records, Hearings, and Legislative Testimony

- Many states, including Utah, place a priority on ensuring access to meetings, records, hearings, and legislative proceedings.
- Although these mechanisms do not always require extensive dialogue or a formal response to public input, they establish that important public business should be conducted in public and be open to scrutiny.

These traditions provide a recognizable foundation on which stronger consultation and accountability requirements can be built.

Citizen Planning and Comprehensive Planning Processes

- Planning commissions, advisory boards, and comprehensive planning processes commonly publish plans and supporting materials, discuss alternatives, solicit feedback, and hold public hearings.
- Officials may be required to make findings or explain why a proposal is consistent with adopted plans and legal standards.

The accountability value lies in the full loop: publication, rationale, alternatives, public feedback, and an explained final decision.

Referenda and Direct Democracy

- Voters may approve or reject statutes, constitutional amendments, taxes, bonds, or other measures depending on state and local law.
- Citizen initiative and referendum processes allow the public to place issues on the ballot or challenge actions taken by elected bodies.

These mechanisms illustrate the principle that some decisions are considered important enough to require direct public authorization or to remain subject to public reversal.

New England Town Meeting

- Proposed actions are publicly announced and supporting information is distributed beforehand.
- Citizens gather in person, hear officials explain and defend proposals, ask questions, challenge assumptions, and propose amendments.
- In some towns, residents vote directly on budgets, ordinances, or other matters.

Its relevance is not limited to direct voting. The deeper principle is that before major community decisions are made, officials must explain what they are proposing, why they are proposing it, and answer questions from the people affected. In New Hampshire and other New England states, this is generally understood not as an ideological innovation but as ordinary local self-government. At the same time, its use has evolved significantly in the past few decades and in many cases diminished.

Public Utility Commission Proceedings

- Utilities seeking rate increases or approval for major investments and infrastructure projects must disclose plans and supporting analyses.
- They are generally required to justify projected costs, demand, benefits, and alternatives in a public record.
- Interested parties may present evidence, question assumptions, object, and seek review; commissions issue public decisions explaining their conclusions.

The intuitive accountability principle is simple: If the public will bear the costs, the public deserves an explanation. The model also demonstrates how technical evidence, public participation, adversarial testing, and a reasoned final decision can be combined.

Bond Measures, Tax Measures, and School Construction Votes

- Communities frequently vote on school bonds, transportation bonds, public facilities, and local tax measures.
- Before the vote, public bodies are expected to explain what is proposed, why it is needed, what it will cost, how it will be financed, and who will pay.

These processes reflect a long-standing expectation that major public expenditures and new financial obligations should be publicly justified and, in specified cases, directly authorized.

Civil Grand Juries

- In some states, civil grand juries or similar grand jury processes are authorized to investigate government operations, public agencies, and public officials.
- These citizen panels may review public records, interview witnesses, examine government practices, and issue public reports with findings and recommendations.
- In some jurisdictions, agencies receiving recommendations are required to provide a formal public response explaining whether they agree, disagree, or plan to implement the recommendations.

These processes illustrate the principle that government actions should remain subject to independent public review and explanation. They also demonstrate that ordinary citizens, when provided with information, time, and a structured process, can evaluate complex public issues and contribute to government accountability.

(Related) Jury Service

- Juries are not a consultation process, but the institution entrusts ordinary citizens with determining facts, weighing evidence, and making consequential decisions.
- Citizens are expected to understand complex material when they receive information, time, procedural protections, and a defined responsibility.

The analogy demonstrates that public involvement need not be limited to expressing uninformed preferences. Citizens can contribute informed judgment when the process is designed to support it.

APPENDIX B

Key Federal Frameworks

National AI Initiative Act (2020)

Established a coordinated federal AI research and policy agenda across agencies.

Executive Order on AI (October 2023, Biden)

Required agencies to assess AI risks, established safety reporting requirements for developers of the most powerful AI systems, and directed federal agencies to develop AI governance frameworks.

Executive Order on AI (December 2025, Trump)

Revoked much of the 2023 order, emphasizing deregulation and AI competitiveness. Directed agencies to develop AI action plans focused on American security and enterprise innovation.

National AI Legislative Framework (March 2026, Trump)

Issued a comprehensive national legislative framework that addresses critical AI policy topics, including safeguarding children, strengthening American communities, respecting intellectual property, free speech protection, ensuring U.S. AI dominance in the world, and preparing for an AI-ready workforce.

Executive Order on Promoting Advanced Artificial Intelligence Innovation and Security (June 2026, Trump)

This order focused on national security, a voluntary review framework for frontier AI models, and cyber defense improvements.

NIST AI Risk Management Framework (2023)

The National Institute of Standards for Technology (NIST) is a voluntary framework for organizations to identify, assess, and manage AI risks. Adopted in government procurement but not widely among industry since the framework is not legally binding or mandatory.

Congressional AI legislation (ongoing)

Multiple AI-related bills have been introduced in Congress, covering transparency, liability, deepfakes, federal worker protections, and AI in hiring. Very few have passed into law as of 2026.

Federal Trade Commission

Has asserted authority over deceptive or unfair AI practices affecting consumers. Has brought actions related to AI-generated content and algorithmic harms.

APPENDIX C

Examples of Legislation in Other States

State approaches to AI regulation continue to evolve rapidly. The examples below illustrate several of the different approaches states have taken to balancing innovation with consumer protection, transparency, accountability, and public safety.

Utah

Utah has pursued a balanced approach that combines support for innovation with targeted consumer protections. **SB 149 (Artificial Intelligence Policy Act, 2024)** created the Office of Artificial Intelligence Policy and established Utah's AI Learning Laboratory (regulatory sandbox). **SB 226 (Artificial Intelligence Amendments, 2025)** narrowed and refined AI disclosure requirements, retained previous enforcement, and created a safe harbor from penalties for businesses who met strict disclosure requirements. AI systems that collect or process personal information are also subject to the **Utah Consumer Privacy Act (UCPA) (SB 227, 2022)**.

California

California has enacted numerous AI-related laws addressing privacy, transparency, election integrity, and AI safety. Among them are the **California Consumer Privacy Act (CCPA) (AB 375, 2018)**, as amended by the **California Privacy Rights Act (CPRA) (Proposition 24, 2020)**, which apply to many AI systems that collect or process personal information. More recently, California enacted **AB 2655 (2024)** and **AB 2839 (2024)** addressing deceptive AI-generated election content, and **SB 942 (2024)** requiring disclosures for certain AI-generated content. California has also considered broader AI safety legislation, including **SB 1047 (2024)**, although that bill was vetoed by Governor Newsom.

Colorado

Colorado became the first state to adopt a comprehensive law governing high-risk AI systems with **SB 24-205 (Consumer Protections for Artificial Intelligence, 2024)**. The law establishes duties for developers and deployers of high-risk AI systems and is intended to reduce algorithmic discrimination in areas such as employment, housing, education, healthcare, insurance, lending, and legal services.

Illinois

Illinois has focused primarily on AI in employment. The **Artificial Intelligence Video Interview Act (HB 2557, 2019)** regulates employers' use of AI to evaluate video interviews and establishes notice, consent, and data retention requirements. The state has subsequently considered and adopted additional measures addressing AI use in employment and consumer protection.

New York

New York has emphasized oversight of AI used in employment. **New York City Local Law 144 of 2021** requires employers using automated employment decision tools to conduct independent bias audits and provide notices to applicants before using such systems. New York has also enacted additional legislation addressing AI-generated election content and other AI-related issues.

Texas

Texas has adopted a targeted approach through legislation addressing specific AI risks rather than broad AI regulation. For example, **HB 149 (2023)** prohibits the creation or distribution of deceptive AI-generated political media within specified periods before an election, while more recent legislation has addressed additional AI-related consumer protection and public safety concerns.

REFERENCES AND RESOURCES

The following sources informed the preparation of this briefing book and are provided for delegates who wish to explore topics in greater depth. Sources are organized by the parts of the briefing book they support.

Part I: Deliberative Democracy & Public Participation

- Fishkin, J. S. (2018). *Democracy when the people are thinking: Reviving our politics through public deliberation*. Oxford University Press. [Foundational text on Deliberative Polling; Stanford Center for Deliberative Democracy: <https://cdd.stanford.edu>]
- OECD. (2020). *Innovative citizen participation and new democratic institutions: Catching the deliberative wave*. OECD Publishing. <https://doi.org/10.1787/339306da-en>
- Utah Legislature. (2024). *Interim study committees overview*. Utah State Legislature. <https://le.utah.gov/asp/interim/Commit.asp>
- University of Utah Hinckley Institute of Politics. (2024). *Public forums and civic engagement programs*. <https://hinckley.utah.edu>

Part II: Understanding Artificial Intelligence

- National Institute of Standards and Technology (NIST). (2023). *Artificial intelligence risk management framework (AI RMF 1.0)*. U.S. Department of Commerce. <https://www.nist.gov/itl/ai-risk-management-framework>
- McKinsey Global Institute. (2023). *The economic potential of generative AI: The next productivity frontier*. McKinsey & Company. <https://www.mckinsey.com/capabilities/mckinsey-digital/our-insights/the-economic-potential-of-generative-ai>
- Acemoglu, D. (2024). The simple macroeconomics of AI. *NBER Working Paper No. 32487*. National Bureau of Economic Research. <https://www.nber.org/papers/w32487> [Examines AI's realistic labor market and productivity effects]
- Goldman Sachs Global Investment Research. (2023). *The potentially large effects of artificial intelligence on economic growth*. Goldman Sachs. [Job displacement and productivity projections from generative AI]

International Energy Agency (IEA). (2024). *Electricity 2024: Analysis and forecast to 2026*. IEA. <https://www.iea.org/reports/electricity-2024> [Data on data center energy consumption growth]

Li, P., Yang, J., Islam, M. A., & Ren, S. (2023). Making AI less “thirsty”: Uncovering and addressing the secret water footprint of AI models. *arXiv preprint arXiv:2304.03271*. <https://arxiv.org/abs/2304.03271> [Water consumption of AI training and inference]

Box Elder County. (n.d.). *Stratos Project Fact Sheet*. Box Elder County, UT. <https://www.boxeldercountyut.gov/647/Stratos-Project-Fact-Sheet>

Rivera, D. (2026). *Utah’s Stratos project draws warnings from Virginia’s “Data Center Alley.”* KSL.com. <https://www.ksl.com/article/51595304/utahs-stratos-project-draws-warnings-from-virginias-data-center-alley>

Additional Data Sources Include:

Per-Query Energy Consumption Table (Gemini, ChatGPT, Claude, Grok)

- Source: Digit.in article “ChatGPT to Gemini: How much energy do your AI queries cost?” which compiled estimates from: ChatGPT (0.3–0.34 Wh per query per Sam Altman), Gemini (0.24 Wh per query from Google’s August 2025 report), Claude (0.3–1.7 Wh per query), Perplexity AI (0.3 Wh), and Grok (2–3 Wh); digit

Claude Model Variations

- Claude 3 Opus (4.05 Wh) and Claude 3 Haiku (0.22 Wh) from an article comparing ChatGPT vs Claude carbon footprints; Carbon Credits
- Claude 3.7 Sonnet (0.84 Wh) from a Devera analysis on AI energy and water usage; Devera

Complex Tasks

- Reasoning models (o3, DeepSeek-R1) consuming 23–39 Wh from an arxiv study “How Hungry is AI? Benchmarking Energy, Water, and Carbon Footprint of LLM Inference;” arXiv
- Claude Code session (41 Wh) from Simon Couch’s analysis of his own Claude Code usage; CNaught

Scale Comparisons (daily/annual)

- ChatGPT daily energy (850 MWh) and water consumption from data center statistics; The Network Installers
- U.S. data center consumption (183 TWh in 2024) from Pew Research Center citing IEA estimates; Pew Research Center

Water Consumption

- 100-word AI prompt water consumption (~519 mL) from University of California research; The Network Installers

Part III: The AI Policy Landscape

- U.S. Congress. (2020). *National Artificial Intelligence Initiative Act of 2020* (Div. E of H.R. 6216, 116th Cong.). <https://www.congress.gov/bill/116th-congress/house-bill/6216>
- The White House. (2023, October). *Executive order on the safe, secure, and trustworthy development and use of artificial intelligence*. <https://www.congress.gov/crs-product/R47843>
- The White House. (2025, December). *Eliminating state law obstruction of national artificial intelligence policy* (Executive Order). <https://www.whitehouse.gov/presidential-actions/2025/12/eliminating-state-law-obstruction-of-national-artificial-intelligence-policy/>
- The White House. (2026, March). *National AI legislative framework*. <https://www.whitehouse.gov/releases/2026/03/president-donald-j-trump-unveils-national-ai-legislative-framework/>
- Utah Legislature. (2024). *Utah Artificial Intelligence Policy Act (UAIPA)*, HB/SB 149. <https://le.utah.gov/~2024/bills/static/SB0149.html>
- Utah Division of Technology Services. (2024). *State agency AI initiatives and inventories*. <https://dts.utah.gov/what-we-are-doing/our-initiatives/artificial-intelligence/>
- Utah Office of Artificial Intelligence Policy (OAIP). (2024). *Office overview and policy resources*. Utah Department of Commerce. <https://commerce.utah.gov/ai/>
- Utah Consumer Privacy Act (UCPA). (2022). S.B. 227, Utah 64th Legislature. <https://commerce.utah.gov/dcp/ucpa/>
- European Parliament and Council of the EU. (2024). *Regulation (EU) 2024/1689 laying down harmonised rules on artificial intelligence (EU AI Act)*. <https://artificialintelligenceact.eu/>
- White & Case LLP. (2024). *AI Watch: Global regulatory tracker — China*. <https://www.whitecase.com/insight-our-thinking/ai-watch-global-regulatory-tracker-china>
- OECD. (2024). *OECD AI principles overview*. <https://oecd.ai/en/ai-principles> [Voluntary principles adopted by 40+ countries]
- Federal Trade Commission. (2024). *Artificial intelligence resources and enforcement actions*. <https://www.ftc.gov/ai>
- Disclose AI. (2024). *State AI disclosure tracker — Utah*. <https://discloseai.net/states/utah>

Part IV: Policy Options, Trade-offs & Values

- National Telecommunications and Information Administration (NTIA). (2024, March). *AI accountability policy report: Earned trust through AI system assurance*. U.S. Department of Commerce. <https://www.ntia.gov/issues/artificial-intelligence/ai-accountability-policy-report> [Recommends independent audits, disclosure standards, and liability rules for high-risk AI systems]
- Future of Privacy Forum (FPF). (2025, October). *The state of state AI: Legislative approaches to AI in 2025*. <https://fpf.org/blog/the-state-of-state-ai-legislative-approaches-to-ai-in-2025/> [Tracks liability frameworks, regulatory sandboxes, and state AI governance trends including Utah’s HB 452]
- U.S. Government Accountability Office (GAO). (2025, April). *Artificial intelligence: Generative AI’s environmental and human effects* (GAO-25-107172). <https://www.gao.gov/products/gao-25-107172> [Policy options for Congress, agencies, and state governments on AI’s energy, water, and labor effects]
- U.S. Government Accountability Office (GAO). (2025, July). *Artificial intelligence: Generative AI use and management at federal agencies* (GAO-25-107653). <https://www.gao.gov/products/gao-25-107653> [Generative AI use cases in government nearly doubled 2023–2024; covers compliance challenges and oversight gaps]
- Congressional Research Service (CRS). (2025). *Regulating artificial intelligence: U.S. and international approaches and considerations for Congress* (R48555). <https://www.congress.gov/crs-product/R48555> [Nonpartisan overview of federal AI bills, state patchwork, and international regulatory frameworks]

Additional Resources for Delegates

- Stanford Human-Centered Artificial Intelligence (HAI). (2025). *Artificial intelligence index report 2025*. <https://hai.stanford.edu/ai-index/2025-ai-index-report> [Comprehensive annual data on AI capabilities, investment, workforce, policy, and public attitudes; notes a widening gap between AI capability and governance readiness]
- Brennan, K., Kak, A., & Myers West, S. (2025, June). *Artificial power: AI Now 2025 landscape report*. AI Now Institute. <https://ainowinstitute.org/publications/research/ai-now-2025-landscape-report> [Examines concentration of power in the AI market; offers strategies for community organizers and policymakers to shape governance]
- Partnership on AI. (2025). *Annual report 2025: Advancing safe and responsible AI systems*. <https://partnershiponai.org/annual-report-2025/> [Multi-stakeholder coalition covering AI agents, synthetic media, global governance, and responsible AI best practices for 2025]
- International Telecommunication Union (ITU). (2025). *The annual AI governance report 2025: Steering the future of AI*. <https://www.itu.int/epublications/publication/the-annual-ai-governance-report-2025-steering-the-future-of-ai> [Global overview of national AI strategies, the rise of AI agents, and multilateral governance frameworks as of mid-2025]
- National Conference of State Legislatures (NCSL). (2024–2026). *Artificial intelligence 2024 legislation*. <https://www.ncsl.org/technology-and-communication/artificial-intelligence-2024-legislation> [Continuously updated tracker of AI-related bills introduced and enacted across all 50 states]
- ITIF Center for Data Innovation. (2025–2026). *AI policy research: Innovation, governance, and competitiveness*. <https://itif.org/centers/data-innovation/> [Pro-innovation perspective on AI regulation, agentic AI policy, public sector AI adoption, and US-China AI competitiveness]

